

MaGNet: A Mamba Dual-Hypergraph Network for Stock Prediction via Temporal-Causal and Global Relational Learning

Peilin Tan

University of California, San Diego
La Jolla, CA, USA
p9tan@ucsd.edu

Dian Tu

University of California, San Diego
La Jolla, CA, USA
ditu@ucsd.edu

Chuanqi Shi

University of California, San Diego
La Jolla, CA, USA
chs028@ucsd.edu

Liang Xie*

Wuhan University of Technology
Wuhan, Hubei, China
whutxl@hotmail.com

Abstract

Stock trend prediction is crucial for profitable trading strategies and portfolio management yet remains challenging due to market volatility, complex temporal dynamics and multifaceted inter-stock relationships. Existing methods struggle to effectively capture temporal dependencies and dynamic inter-stock interactions, often neglecting cross-sectional market influences, relying on static correlations, employing uniform treatments of nodes and edges, and conflating diverse relationships. This work introduces MaGNet, a novel **Mamba dual-hyperGraph Network** for stock prediction, integrating three key innovations: (1) a MAGE block, which leverages bidirectional Mamba with adaptive gating mechanisms for contextual temporal modeling and integrates a sparse Mixture-of-Experts layer to enable dynamic adaptation to diverse market conditions, alongside multi-head attention for capturing global dependencies; (2) Feature-wise and Stock-wise 2D Spatiotemporal Attention modules enable precise fusion of multivariate features and cross-stock dependencies, effectively enhancing informativeness while preserving intrinsic data structures, bridging temporal modeling with relational reasoning; and (3) a dual hypergraph framework consisting of the Temporal-Causal Hypergraph (TCH) that captures fine-grained causal dependencies with temporal constraints, and Global Probabilistic Hypergraph (GPH) that models market-wide patterns through soft hyperedge assignments and Jensen-Shannon Divergence weighting mechanism, jointly disentangling localized temporal influences from instantaneous global structures for multi-scale relational learning. Extensive experiments on six major stock indices demonstrate MaGNet outperforms state-of-the-art methods in both superior predictive performance and exceptional investment returns with robust risk management capabilities. **Datasets, source code, and model weights are available at our GitHub repository: <https://github.com/PeilinTime/MaGNet>.**

Keywords

Stock Prediction, Mamba, Hypergraph Neural Network

1 Introduction

The stock market serves as a crucial component of the global financial system and a primary investment avenue. Stock trend prediction, which forecasts future price movements to guide investment

decisions and risk management, has garnered substantial attention. However, prediction remains inherently challenging due to high volatility, non-stationary behavior and complex influencing factors including macroeconomic conditions, company performance, and inter-stock relationships.

Stock price prediction has evolved from traditional methods (SVM, ARIMA [1]) to deep learning approaches. While CNNs [19], RNNs (LSTMs [4]), and Transformers [22] improved capturing complex market behaviors, they face challenges with long-range dependencies and computational complexity. State Space Models like Mamba [16] recently emerged as efficient alternatives, achieving near-linear complexity through selective scan mechanisms.

Stock correlation modeling progressed from static, predefined connections to dynamic representations. Graph Neural Networks (GNNs) [5] model stocks as nodes with dynamic edges but only capture pairwise relationships. Since stock movements often involve higher-order group dynamics through shared industry membership or ownership, Hypergraph Neural Networks (HGNNs) [9] were introduced to connect multiple nodes simultaneously via hyperedges, though challenges remain in appropriately weighting neighbors and hyperedges.

Despite these advances, critical limitations persist in current approaches that need addressing: (1) For time series modeling, while Mamba offers efficient linear complexity through selective state space mechanisms, it lacks contextual modeling with sophisticated temporal fusion, adaptation to diverse market regimes, and global dependency capture—all crucial for financial markets. Moreover, existing temporal models treat each stock's time series independently, failing to preserve cross-sectional information embedded in the feature space that captures critical inter-stock dynamics [24]; (2) For relational modeling, while HGNNs advance beyond pairwise connections, they suffer from uniform treatment of nodes within hyperedges despite varying influence levels, and inability to distinguish between different types of relationships (e.g., causal vs. instantaneous, local vs. global). These limitations result in models that cannot fully capture the complex, dynamic and multi-scale nature of financial markets, necessitating a more sophisticated architecture.

To address these challenges, we propose a novel architecture combining advanced temporal modeling with dynamic relational learning. Our contributions are threefold:

*Corresponding author.

- **MAGE Block with 2D Attention:** To fully capture temporal dynamics, we design the **MAGE** (Mamba-Attention-Gating-Experts) block by enhancing bidirectional Mamba with adaptive gating to capture the full temporal context. We integrate sparse MoE for market regime adaptation and multi-head attention for global dependencies. We further apply feature-wise 2D spatiotemporal attention to enhance stock features by capturing dynamic interactions across features and stocks over time. This design jointly captures temporal dynamics and spatial patterns.
- **Dual Hypergraph Framework:** To model dynamic and high-order relations among stocks, we introduce two complementary hypergraphs. The Temporal-Causal Hypergraph (TCH) captures fine-grained, localized relations across stocks and time. The Global Probabilistic Hypergraph (GPH) encodes broader market structures through probabilistic hyperedges, allowing stocks to have varying membership degrees across multiple groups rather than uniform treatment. This dual design disentangles local temporal-causal signals from global market patterns, enabling expressive and flexible relational learning.
- **State-of-the-Art Performance:** Extensive experiments on six major stock indices demonstrate that MaGNet achieves superior predictive accuracy (up to 54.9% on CSI 300) and exceptional investment returns with robust risk-adjusted performance, including Sharpe ratios exceeding 1.0 on multiple markets and annual returns up to 22.6%.

2 Related Work

2.1 Time Series Models for Stock Prediction

Deep learning revolutionized stock prediction by capturing complex temporal patterns. RNNs and variants (LSTM, GRU) became widely adopted for modeling sequential dependencies [33]. CNN-based approaches treated historical data as feature maps, with architectures like dilated CNNs capturing multi-scale patterns [34]. Attention mechanisms enhanced these models by focusing on relevant historical states [30], culminating in transformer architectures that achieved remarkable results through self-attention [8].

Recently, State Space Models (SSMs) emerged as efficient alternatives, combining CNN-like parallel training with RNN-like fast inference. Mamba [16], a selective SSM with time-varying parameters and hardware-aware algorithms, has shown particular promise. Extensions include S-Mamba [39] for multivariate series and bidirectional variants like MambaMixer [3] for enhanced contextual modeling. However, sequential models often treat stocks independently, ignoring crucial inter-stock relationships.

2.2 Graph and Hypergraph Methods for Inter-Stock Relationships

Recognizing stock interdependencies, graph-based methods explicitly model relationships through various connections: shareholding, industry sectors [18], supply chains, and price correlations [27]. Graph Neural Networks (GNNs) learn representations over graphs through neighborhood aggregation, with GCNs combined with temporal models [12] capturing evolving relationships and GATs [25, 36] using attention for refined message passing.

However, GNNs face limitations: reliance on predefined relationships that may be incomplete or noisy, and inability to model higher-order group interactions [13, 44]. Hypergraph methods [7, 14, 41] address this by using hyperedges connecting multiple nodes simultaneously, naturally representing group-wise dependencies like industry sectors or shared fund ownership. Frameworks include HyperGCN [42] and STHGCN [31] for spatiotemporal modeling. Despite these advances, many hypergraph models struggle with dynamic temporal features and uniform node treatment.

3 Methodology

3.1 Problem Definition

We address the task of predicting next-day stock price movement, framed as a binary classification problem. Let $\mathcal{S} = \{s_1, s_2, \dots, s_N\}$ represent a set of N stocks in the market. For each stock $s_i \in \mathcal{S}$, we consider its historical trading data over a lookback window of T days, with F financial indicators recorded daily, resulting in a feature matrix $\mathbf{X}_i \in \mathbb{R}^{T \times F}$ for stock s_i . The complete input data across all stocks can be represented as $\mathbf{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_N] \in \mathbb{R}^{N \times T \times F}$. The prediction target is the direction of all stocks' closing price movement on the following day, encoded as a binary label:

$$y_i = \begin{cases} 1, & \text{if } p_i^{(t+1)} > p_i^{(t)}, \\ 0, & \text{otherwise} \end{cases}, \forall i \in \{1, 2, \dots, N\}, \quad (1)$$

where $p_i^{(t)}$ denotes the closing price of stock s_i on day t .

The objective is to learn a predictive function $f(\cdot; \Theta)$ parameterized by Θ that maps \mathbf{X} to the predicted probabilities of upward price movement for all stocks: $\hat{\mathbf{Y}} = f(\mathbf{X}; \Theta)$, $\hat{\mathbf{Y}} \in [0, 1]^N$.

3.2 Feature Embedding

To compress information and extract salient features, each stock's daily feature vector $\mathbf{x}_{n,t} \in \mathbb{R}^F$ is first projected into a shared latent space through a feed-forward embedding network:

$$\mathbf{z}_{n,t} = \text{Embedding}(\mathbf{x}_{n,t}) \in \mathbb{R}^D. \quad (2)$$

3.3 MAGE Block

Figure 1 illustrates the overall architecture of MaGNet, with the **MAGE** (Mamba-Attention-Gating-Experts) block serving as its core temporal modeling component. The MAGE block processes the embedded representations to capture heterogeneous temporal interactions through four synergistic modules: bidirectional Mamba first captures forward and backward temporal patterns, followed by a gating mechanism that adaptively fuses these contextual representations; then passes through a sparse Mixture of Experts layer for specialized processing under different market regimes; and finally, multi-head self-attention incorporates global dependencies across the entire sequence. This design balances computational efficiency with modeling capacity for complex market dynamics.

3.3.1 Bidirectional Mamba. Mamba [16] enhances SSMs with selective scan mechanism and data-dependent parameters, achieving near-linear complexity while modeling long-range dependencies without requiring positional encodings [29]. However, standard Mamba processes sequences unidirectionally, limiting contextual

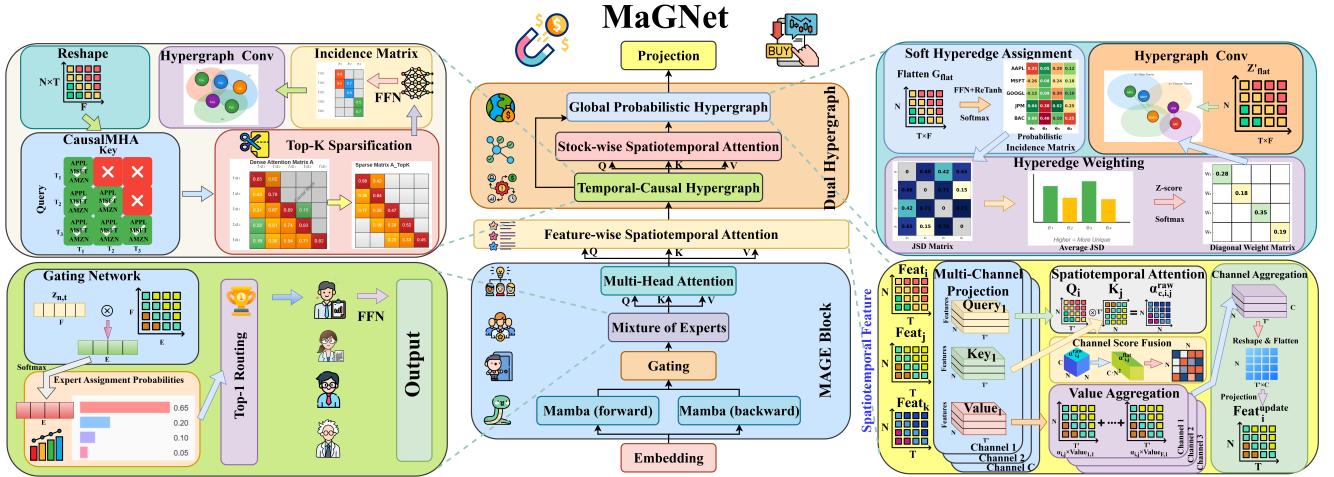


Figure 1: Overview of MaGNet architecture. It consists of: (1) MAGE Block, combining bidirectional Mamba, adaptive gating, MoE, and multi-head attention for temporal modeling; (2) Feature-wise 2D Spatiotemporal Attention to capture cross-feature dependencies while preserving spatiotemporal structure; (3) Dual Hypergraph Module, with TCH modeling fine-grained temporal-causal relations and GPH capturing global market patterns via soft assignments and JSD-based weighting.

understanding. We employ bidirectional Mamba [2] to address this:

$$Z_n^{\text{fwd}} = \text{Mamba}(Z_n), \quad (3)$$

$$Z_n^{\text{bwd}} = \text{Reverse}(\text{Mamba}(\text{Reverse}(Z_n))). \quad (4)$$

where $Z_n = [z_{n,1}, z_{n,2}, \dots, z_{n,T}] \in \mathbb{R}^{T \times D}$ and $\text{Reverse}(\cdot)$ denotes time-axis reversal. This bidirectional processing captures both forward and backward temporal dependencies, providing more comprehensive sequence representations essential for complex financial patterns.

3.3.2 Gating Mechanism. To fuse the forward and backward representations Z_n^{fwd} and Z_n^{bwd} into a unified encoding while avoiding naive averaging or concatenation, we employ a Gating Mechanism [6] that adaptively integrates these two directional outputs. At time step t :

$$z_t^G = \text{Gate}(z_t^{\text{fwd}}, z_t^{\text{bwd}}) \quad (5)$$

$$= \sigma(W_f z_t^{\text{fwd}} + b_f + W_b z_t^{\text{bwd}} + b_b), \quad (6)$$

where $W_f, W_b \in \mathbb{R}^{D \times D}$, $b_f, b_b \in \mathbb{R}^D$ are learnable parameters. This Gating Mechanism adaptively controls how much information to retain from each direction, enabling context-aware fusion for better temporal representation.

3.3.3 Mixture of Experts. Characterized by intrinsic volatility and structural heterogeneity, stock market environments pose significant challenges for distributionally robust modeling. Training data may originate from bullish markets, whereas evaluation could occur under bearish conditions or amid abrupt distributional shifts—such as major policy changes or global events. To address this, we incorporate a Switched Mixture-of-Experts (MoEs) [32] layer with E experts that account for a wide range of possible market scenarios, enabling sparse conditional computation and enhancing model resilience without incurring scaling cost. First, the gating network

computes expert assignment probabilities:

$$\mathbf{P} = \text{Softmax}(\mathbf{Z}^G \mathbf{W}_g) \in \mathbb{R}^{N \times T \times E}, \quad (7)$$

where $\mathbf{W}_g \in \mathbb{R}^{D \times E}$. We employ Top-1 routing using capacity-based normalization with scaling factor C to enforce sparse expert selection and balanced utilization:

$$\hat{p}_{n,t,e} = \begin{cases} p_{n,t,e}, & \text{if } e = \arg \max_j p_{n,t,j}, \\ 0, & \text{otherwise} \end{cases}, \quad (8)$$

$$\tilde{p}_{n,t,e} = C \cdot \frac{\hat{p}_{n,t,e}}{\sum_{n',t'} \hat{p}_{n',t',e}}. \quad (9)$$

Each expert \mathcal{E}_e is a position-wise feed-forward network:

$$\mathbf{z}_{n,t}^{\text{MoE}} = \mathcal{E}_e(z_{n,t}^G), \quad e^* = \arg \max_j p_{n,t,j}, \quad (10)$$

$$\mathcal{E}_e(\mathbf{z}) = \mathbf{W}_e^{(2)} \text{GELU}(\mathbf{W}_e^{(1)} \mathbf{z} + \mathbf{b}_e^{(1)}) + \mathbf{b}_e^{(2)}. \quad (11)$$

This design enables dynamic capacity allocation across time steps while maintaining load balancing through the capacity normalization mechanism.

3.3.4 Multi-Head Self-Attention. Finally, to complement bidirectional representations with global dependencies across the entire sequence, we incorporate multi-head attention [35]. Given input $\mathbf{Z}^{\text{MoE}} \in \mathbb{R}^{N \times T \times D}$, we compute h attention heads in parallel across multiple representation subspaces with size $d_h = D/h$ and combine them:

$$\mathbf{Q}_i = \mathbf{Z}^{\text{MoE}} \mathbf{W}_i^Q, \quad \mathbf{K}_i = \mathbf{Z}^{\text{MoE}} \mathbf{W}_i^K, \quad \mathbf{V}_i = \mathbf{Z}^{\text{MoE}} \mathbf{W}_i^V, \quad (12)$$

$$\text{head}_i = \text{Softmax}\left(\frac{\mathbf{Q}_i \mathbf{K}_i^\top}{\sqrt{d_k}}\right) \mathbf{V}_i, \quad (13)$$

$$\mathbf{Z}^{\text{MAGE}} = \text{Concat}(\text{head}_1, \dots, \text{head}_h) \mathbf{W}^O, \quad (14)$$

where $\mathbf{W}_i^Q, \mathbf{W}_i^K, \mathbf{W}_i^V \in \mathbb{R}^{D \times d_h}, \forall i \in \{1, \dots, h\}$ and $\mathbf{W}^O \in \mathbb{R}^{D \times D}$. This module enables the model to jointly attend to information

from different representation subspaces, effectively capturing global dependencies.

Together, these four components form a flexible and powerful architecture capable of learning diverse temporal patterns in financial time series.

3.4 Feature-wise 2D Spatiotemporal Attention

Traditional time series methods focus on single dimension (typically temporal), requiring multivariate data to be flattened or processed separately, disrupting spatiotemporal structure and may obscure meaningful relationships across time and assets. While iTransformer [26] models cross-feature dependencies, it overlooks the crucial spatiotemporal relationships. We propose Feature-wise 2D Spatiotemporal Attention that preserves this structure by representing each feature as an $N \times T$ matrix (stocks \times time), enabling direct feature-to-feature interactions while maintaining spatiotemporal integrity, yielding rich representations for downstream dual hypergraph learning.

The module first transposes $Z^{\text{MAGE}} \in \mathbb{R}^{N \times T \times D}$ into $Z' \in \mathbb{R}^{D \times N \times T}$, where each feature d is represented by matrix $Z'_d \in \mathbb{R}^{N \times T}$, capturing dynamics across all stocks and time steps. To handle the increased element dimensionality while maintaining computational efficiency, we employ C parallel attention channels. For each feature d and channel $c \in \{1, \dots, C\}$, we compute query, key and value projections: $Q_{d,c} = Z'_d W_c^Q + b_c^Q$, $K_{d,c} = Z'_d W_c^K + b_c^K$, $V_{d,c} = Z'_d W_c^V + b_c^V$, where $W_c^Q, W_c^K, W_c^V \in \mathbb{R}^{T \times T'}$ and $b_c^Q, b_c^K, b_c^V \in \mathbb{R}^{1 \times T'}$ are learnable parameters. To capture feature-to-feature relationships, we compute cross-stock multi-channel attention scores between each feature pair $i, j \in \{1, \dots, D\}$, and then fuse them through a feed-forward network:

$$\alpha_{i,j}^{\text{raw}} = \frac{1}{\sqrt{T'}} \left[Q_{i,1} K_{j,1}^\top, \dots, Q_{i,C} K_{j,C}^\top \right] \in \mathbb{R}^{C \times N \times N}, \quad (15)$$

$$\alpha_{i,j}^{\text{flat}} = \text{Flatten}(\alpha_{i,j}^{\text{raw}}) \in \mathbb{R}^{C \cdot N^2}, \alpha'_{i,j} = \text{FFN}(\alpha_{i,j}^{\text{flat}}) \in \mathbb{R}. \quad (16)$$

The aggregated representations are computed via weighted summation:

$$\mathbf{B}'_{i,c} = \sum_{j=1}^D a_{i,j} V_{j,c} \in \mathbb{R}^{N \times T'}, \quad a_{i,j} = \text{softmax}(\alpha'_{i,j}). \quad (17)$$

All channels' outputs $\{\mathbf{B}'_{i,c}\}_{c=1}^C$ are then stacked into $\mathbf{B}'_i \in \mathbb{R}^{C \times N \times T'}$, reshaped to $\mathbf{B}_i \in \mathbb{R}^{N \times (T' \cdot C)}$, and projected to dimension T via feed-forward network to obtain $\mathbf{z}_i^{\text{F2D}} \in \mathbb{R}^{N \times T}$. Stacking all features yields $\mathbf{Z}^{\text{F2D}} \in \mathbb{R}^{N \times T \times D}$ with enriched cross-feature dependencies.

3.5 Dual-Hypergraph Learning

Traditional graph methods capture only pairwise relationships, missing the group dynamics where multiple stocks move synchronously. Hypergraphs address this by connecting multiple nodes via hyperedges. We propose a dual hypergraph framework leveraging the spatiotemporal representations from our 2D attention mechanisms: the Temporal-Causal Hypergraph (TCH) learns dynamic causal relationships respecting temporal ordering, while the Global Probabilistic Hypergraph (GPH) captures instantaneous market-wide patterns through soft hyperedge assignments. This design models both localized temporal influences and global market structures for multi-scale market interdependencies.

3.5.1 Temporal-Causal Hypergraph (TCH). The Temporal-Causal Hypergraph (TCH) discovers high-order causal dependencies among stocks at the fine-grained stock-time level—where each stock at a specific timestamp can influence multiple others at current or subsequent timestamps. Unlike static graph approaches, TCH adaptively learns dynamic hypergraph topology that respects temporal ordering, tracking how stock movements propagate through the market over time. We restructure the input \mathbf{Z}^{F2D} by flattening:

$$\mathbf{Z}_{\text{flat}} = \text{Flatten}((\mathbf{Z}^{\text{F2D}})^\top) \in \mathbb{R}^{(T \cdot N) \times D}, \quad (18)$$

where each row represents a time-stock pair. We then apply Causal Multi-Head Attention (CausalMHA) to capture causal relationships while preventing future information leakage:

$$\mathbf{A} = \text{CausalMHA}(\mathbf{Z}_{\text{flat}}) \in \mathbb{R}^{(T \cdot N) \times (T \cdot N)}. \quad (19)$$

The causality constraint is enforced through a upper triangular block mask such that each node can only attend to others from current or earlier time steps. This explicit design captures temporal and cross-sectional dependencies with causal validity, modeling how past conditions and peer movements influence future stock behavior.

The dense attention matrix \mathbf{A} may contain weak or noisy connections that can impair performance. To extract the most salient relationships, we apply Top- K sparsification:

$$\mathbf{A}_{\text{TopK}}[i, j] = \begin{cases} \text{Softmax}(\mathbf{A}[i, :])_j, & \text{if } j \in \text{Top-}K(i) \\ 0, & \text{otherwise} \end{cases}. \quad (20)$$

This sparsification highlights the strongest causal paths, reduces computational complexity, and mitigates the influence of noisy or insignificant connections.

To transform pairwise attention scores into the hypergraph structure, we introduce a FFN with a novel Rectified Tanh (ReTanh) activation:

$$\mathbf{Z}_1 = \text{ReTanh}(\mathbf{W}_1 \mathbf{A}_{\text{TopK}}), \mathbf{H}_{\text{TCH}} = \text{ReTanh}(\mathbf{W}_2 \mathbf{Z}_1), \quad (21)$$

where $\mathbf{H}_{\text{TCH}} \in \mathbb{R}^{(T \cdot N) \times M_1}$ is incidence matrix, M_1 is the number of hyperedges. The ReTanh activation is defined as:

$$\text{ReTanh}(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ \tanh(x) & \text{if } x > 0 \end{cases}. \quad (22)$$

ReTanh enhances hypergraph learning by combining sparsity through rectified filtering of weak linkages with stability from tanh-bounded hyperedge assignments, reducing outlier impact like abnormal financial events, and ensuring robust, balanced correlations in noisy market environments.

Next, we perform hypergraph convolution to enable high-order information propagation:

$$\mathbf{Z}'_{\text{flat}} = \text{ELU}(\mathbf{H}_{\text{TCH}} \mathbf{H}_{\text{TCH}}^\top \mathbf{Z}_{\text{flat}} \mathbf{P}_1), \quad (23)$$

where $\mathbf{P}_1 \in \mathbb{R}^{D \times D}$ is a learnable projection matrix. This operation enables each node to aggregate information from all nodes within its hyperedges, capturing complex group dynamics that pairwise methods miss. Finally, we reshape the transformed features back to the original format $\mathbf{Z}^{\text{TCH}} = \text{Reshape}(\mathbf{Z}'_{\text{flat}}) \in \mathbb{R}^{N \times T \times D}$.

TCH achieves adaptive learning of temporal-causal relationships that respect market dynamics. The combination of causal attention for temporal validity and hypergraph convolution for high-order

modeling creates a powerful framework for capturing the complex interdependencies that drive stock market movements.

3.5.2 Global Probabilistic Hypergraph (GPH). While TCH captures temporal-causal relationships, financial markets also exhibit instantaneous global patterns. The Global Probabilistic Hypergraph (GPH) discovers these market-wide group interactions through probabilistic hyperedges, allowing stocks to participate in multiple market themes simultaneously with varying membership degrees.

To model direct inter-stock relationships, we first introduce Stock-wise 2D Spatiotemporal Attention, applying similar mechanism as its feature-wise counterpart but along the stock dimension, where each stock n is represented by its temporal features $Z_n^{\text{TCH}} \in \mathbb{R}^{T \times D}$, producing output $Z^{\text{N2D}} \in \mathbb{R}^{N \times T \times D}$ that encodes cross-stock dependencies.

To learn soft hyperedge assignments, we flatten Z^{N2D} into $G_{\text{flat}} \in \mathbb{R}^{N \times (T \cdot D)}$, apply a FFN with ReTanh activation, and normalize column-wise to obtain the probabilistic incidence matrix $H_{\text{GPH}} \in \mathbb{R}^{N \times M_2}$:

$$H_{\text{GPH}} = \text{softmax}(\text{ReTanh}(\text{FFN}(G_{\text{flat}}))), \quad (24)$$

where M_2 is the number of hyperedges. Each column e_j in H_{GPH} defines a soft hyperedge as a probability distribution with membership probabilities summing to 1.

To address redundant relationships between hyperedges, we weight each hyperedge by its distinctiveness using Jensen-Shannon Divergence (JSD). For each hyperedge, we compute its average divergence:

$$\mu_j = \frac{1}{M_2} \sum_{i=1}^{M_2} \text{JSD}(e_i \| e_j), \quad (25)$$

JSD's symmetry and boundedness ($[0, \log 2]$) ensure fair evaluation regardless of comparison order and stable optimization. The importance weight is:

$$w_j = \text{softmax}(\text{Z-score}(\mu_j)), \quad (26)$$

where Z-score prevents domination by only a small number of hyperedges. This JSD weighting scheme assigns higher weights to unique hyperedges that capture distinct and informative market structures.

Finally, the global hypergraph convolution integrates the weighted hyperedges to propagate information across all stocks:

$$Z^{\text{GPH}} = \text{ELU}(H_{\text{GPH}} W H_{\text{GPH}}^T Z'_{\text{flat}} P_2) \in \mathbb{R}^{N \times (T \cdot D)}, \quad (27)$$

where $W = \text{diag}(w_1, \dots, w_{M_2})$ is diagonal weight matrix and $P_2 \in \mathbb{R}^{(T \cdot D) \times (T \cdot D)}$ is a learnable projection matrix.

GPH propagates information globally via weighted group memberships and importance scores. Combined with TCH's temporal-causal patterns, this dual hypergraph framework captures both localized temporal influences and global market structures, providing the multi-scale representational capacity essential for accurate stock prediction.

4 Experiments

4.1 Datasets

We evaluated our method on six major stock indices (DJIA, HSI, NASDAQ 100, S&P 100, CSI 300, and Nikkei 225) using data from January 1, 2020 through December 31, 2024. Data was split 7:1:2 for

training, validation (hyperparameter selection) and testing (evaluation). See Table 1 for dataset statistics.

Table 1: Statistics of Datasets

Dataset	# Stocks	# Training	# Val	# Test
DJIA	30	879	125	253
HSI	71	860	122	247
NASDAQ 100	92	879	125	253
S&P 100	99	879	125	253
CSI 300	215	848	121	243
Nikkei 225	222	855	122	245

4.2 Features

We obtained historical stock data from Yahoo Finance ¹, collecting five attributes: close, high, low, open, volume. To enhance features, we used Qlib [43] to compute Alpha158 and Alpha360 technical indicators. After filtering out missing values, we combined these features with the five attributes to create an enriched dataset. We applied Z-Score normalization independently to each data split to prevent information leakage while maintaining stable training process.

4.3 Baselines

We evaluate our method against 17 baselines spanning three categories:

- **Stock Prediction Models** (6): SFM [46], Adv-ALSTM [11], DTML [44], ESTIMATE [20], StockMixer [10], MASTER [23];
- **Time Series Models** (8): GRU [6], LSTM [15], DLinear [45], TimesNet [40], PatchTST [28], iTransformer [26], TimeMixer [37], TimeXer [38];
- **Graph Models** (3): GCN [21], GraphSAGE [17], GAT [36].

Baselines' descriptions are provided in Appendix A.1.

4.4 Evaluation Metrics

We evaluate model's performance using classification metrics and portfolio backtesting. For predictive abilities, we use Accuracy (ACC), Precision (PRE), Recall (REC), F1 score, and AUC. To assess model's profitability and risk in simulated investment scenarios, we employ Annual Return (AR), Sharpe Ratio (SR, applying a 2% risk-free rate), Calmar Ratio (CR), and Maximum Drawdown (MDD). Detailed definitions and formulas of all metrics are provided in Appendix A.2.

4.5 Hyperparameter Settings

We passed the final representations Z^{GPH} through a feed-forward network followed by a softmax layer to predict the probability distribution of next-day stock movement directions. The model was implemented in PyTorch and optimized with cross-entropy loss. Hyperparameters were selected via grid search to maximize

¹<https://ranaroussi.github.io/yfinance/>

Table 2: Prediction performance comparison on DJIA, HSI, and NASDAQ 100. The best results are in bold and the second-best results are underlined.

Model	DJIA					HSI					NASDAQ 100				
	ACC	PRE	REC	F1	AUC	ACC	PRE	REC	F1	AUC	ACC	PRE	REC	F1	AUC
GRU	51.59	53.13	75.91	62.51	49.98	51.89	49.12	1.73	3.34	50.78	52.07	52.37	88.86	65.90	51.41
LSTM	53.03	53.15	98.17	68.97	50.57	51.96	50.35	7.11	12.46	50.75	51.86	52.09	95.07	67.30	50.52
DLinear	52.39	52.98	92.81	67.46	49.65	52.20	51.46	10.25	17.09	50.78	52.18	52.25	95.58	67.57	51.32
TimesNet	50.34	52.53	68.43	59.44	49.04	52.40	50.92	27.99	36.13	52.07	51.77	52.18	89.32	65.88	50.37
PatchTST	49.66	52.62	53.15	52.89	49.04	50.25	47.44	32.13	38.31	49.40	51.72	52.23	86.44	65.11	49.99
iTransformer	<u>53.13</u>	<u>53.39</u>	93.30	67.91	50.11	52.50	51.22	25.49	34.03	52.32	51.09	51.76	90.76	65.92	48.57
TimeM.	52.39	53.02	91.55	67.15	49.36	51.89	49.87	11.42	18.58	51.13	<u>52.64</u>	52.53	94.88	67.62	51.93
TimeXer	52.92	53.07	98.71	69.03	49.06	52.03	50.90	6.65	11.76	50.84	52.50	52.43	<u>95.69</u>	<u>67.74</u>	<u>52.12</u>
GCN	52.30	53.05	89.50	66.61	49.65	51.83	49.81	24.92	33.22	51.52	52.00	<u>52.91</u>	71.86	60.95	50.86
GraphSAGE	52.88	53.03	<u>99.25</u>	<u>69.13</u>	49.59	<u>52.76</u>	51.40	32.13	39.55	52.87	51.47	52.19	82.28	63.86	51.93
GAT	52.90	53.11	97.39	68.74	50.03	52.32	51.39	15.56	23.89	<u>52.93</u>	52.37	52.41	93.86	67.26	50.05
SFM	49.89	53.44	55.02	54.22	48.96	51.11	48.51	<u>33.84</u>	<u>39.87</u>	50.68	51.05	52.82	61.45	56.81	50.81
Adv-ALSTM	52.69	53.07	95.06	68.11	50.72	51.07	48.06	21.94	30.08	49.92	51.90	52.14	94.61	67.21	49.83
DTML	52.30	53.02	90.37	66.83	<u>50.85</u>	51.73	49.69	31.91	38.87	50.82	52.01	52.52	82.49	64.18	51.05
ESTIMATE	52.93	53.15	96.77	68.61	50.82	52.08	50.95	8.91	15.17	50.64	51.94	52.15	94.41	67.19	50.57
StockMixer	52.49	53.10	90.90	67.04	49.16	51.92	50.00	15.59	23.76	50.84	51.70	52.02	94.68	67.14	49.23
MASTER	52.15	53.18	83.53	64.98	49.78	52.58	51.88	19.14	27.97	52.49	51.81	52.38	83.04	64.24	51.92
MaGNet	53.16	53.16	100.00	69.42	51.10	54.19	52.85	43.90	47.96	54.12	53.72	53.09	96.14	68.41	52.24

validation accuracy. Detailed configurations and dataset-specific settings are provided in Appendix A.3.

4.6 Prediction Results

The prediction results on DJIA, HSI and NASDAQ 100 indices are shown in Table 2, and results on S&P 100, CSI300 and Nikkei 255 are presented in Appendix B.1. MaGNet consistently outperforms all baselines across all indices in terms of accuracy, recall, F1-score, and AUC, demonstrating its superior predictive capability. Notably, it achieves the highest accuracy on HSI (54.19%) and competitive performance on DJIA and NASDAQ 100. The model also delivers substantial improvements in recall, particularly for DJIA, indicating strong sensitivity to upward movements. Across these datasets, MaGNet maintains balanced precision–recall trade-offs, reflected in leading or near-leading F1-scores. Its consistently higher AUC values further suggest enhanced discrimination ability compared to both time-series-only and graph-based baselines.

4.7 Backtesting Results

To evaluate the practical profitability of the proposed model in realistic trading scenarios, we conduct daily backtests using a systematic portfolio-based trading strategy that simulates real-world trading mechanics. The detailed description of the trading strategy and the configurations can be found in Appendix A.4.

The backtesting results on DJIA, HSI, and NASDAQ 100 indices are summarized in Table 3, and results on the other three indices are provided in Appendix B.2. Across these indices, MaGNet consistently converts predictive gains into higher risk-adjusted returns, achieving the highest or near-highest Sharpe Ratios and Annual Returns among all models. Drawdowns remain moderate compared

with baselines, reflecting effective downside control. These findings demonstrate that integrating MAGE’s temporal modeling with the dual-hypergraph framework yields robust and profitable trading performance across various indices.

Figure 2 and 3 present the backtesting performance visualization of MaGNet across all six stock indices, with each subplot containing three panels: portfolio value over time (left), daily returns distribution (center), and drawdown trajectory (right). The portfolio value trajectories demonstrate MaGNet’s consistent ability to generate positive returns, with particularly strong performance on CSI300 (achieving approximately 22% growth), DJIA (20% growth), and Nikkei225 (18% growth). The daily returns distributions exhibit near-normal characteristics centered slightly above zero (mean daily return of 0.07–0.09%), indicating the model’s capability to maintain positive expected returns while avoiding extreme tail events. Notably, the drawdown profiles reveal MaGNet’s robust risk management, with maximum drawdowns contained below 5% for CSI300 and DJIA, and remaining under 10% across all indices. The relatively shallow and quickly-recovering drawdown patterns, particularly evident in the DJIA and CSI300 results, corroborate the superior Sharpe ratios reported in Table 3 and 8, demonstrating that MaGNet not only achieves strong absolute returns but does so with controlled downside risk, a critical requirement for practical trading applications.

4.8 Ablation Studies

To evaluate the contribution of each component, we conducted ablation studies by removing key modules from MaGNet. The complete ablation results across all six indices are provided in Appendix B.3. The results on NASDAQ100 are shown in Table 4. The results

Table 3: Backtesting performance on DJIA, HSI, and NASDAQ 100. The best results are in bold and the second-best results are underlined.

Model	DJIA				HSI				NASDAQ 100			
	AR	SR	CR	MDD	AR	SR	CR	MDD	AR	SR	CR	MDD
GRU	18.84	0.62	1.21	15.53	8.38	0.54	0.71	11.85	15.98	0.83	<u>1.79</u>	8.92
LSTM	17.00	1.41	3.30	5.15	5.32	0.25	0.44	11.97	11.68	0.65	<u>1.16</u>	10.10
DLinear	8.81	0.64	1.45	6.07	7.49	0.48	1.02	7.38	10.26	0.56	1.04	9.83
TimesNet	1.99	-0.01	7.95	0.25	6.43	0.50	1.18	5.45	6.36	0.29	0.59	10.83
PatchTST	1.28	-0.22	0.63	<u>2.04</u>	5.31	0.70	2.69	<u>1.97</u>	-3.08	-0.35	-0.24	12.94
iTransformer	9.82	0.74	1.83	5.37	5.45	0.40	1.00	5.45	2.16	0.01	0.20	10.70
TimeMixer	9.80	0.74	1.68	5.82	7.02	0.51	0.94	7.49	12.30	0.70	1.20	10.21
TimeXer	17.68	1.48	3.43	5.16	7.53	0.21	0.40	18.65	11.90	0.66	1.18	10.10
GCN	11.99	0.88	1.71	7.00	8.25	0.43	0.80	10.37	7.47	0.43	0.87	8.60
GraphSAGE	18.29	<u>1.55</u>	3.35	5.46	9.91	0.56	0.95	10.41	4.72	0.23	0.55	8.53
GAT	6.38	0.40	1.09	5.83	4.28	0.16	0.43	9.85	10.65	0.59	1.14	9.37
SFM	6.88	0.70	1.73	3.98	2.09	0.02	<u>2.02</u>	1.04	2.85	0.14	0.67	4.23
Adv-ALSTM	15.01	1.23	2.92	5.14	3.15	0.10	0.30	10.38	11.03	0.62	1.29	8.55
DTML	13.82	1.11	2.71	5.09	8.11	0.35	0.68	11.94	7.11	0.34	0.70	10.19
ESTIMATE	24.40	1.42	3.38	7.21	16.26	<u>0.66</u>	1.01	16.12	<u>16.77</u>	<u>1.03</u>	1.78	9.44
StockMixer	8.46	0.61	1.47	5.76	4.74	0.30	1.00	4.75	10.02	0.54	1.00	9.97
MASTER	3.40	0.13	0.52	6.50	5.19	0.20	0.46	11.33	7.03	0.44	0.83	8.43
MaGNet	<u>19.92</u>	1.70	<u>3.93</u>	5.07	<u>12.25</u>	<u>0.66</u>	1.33	9.20	17.09	1.05	2.09	<u>8.18</u>

Table 4: Ablation results on NASDAQ 100. The best results are in bold and the second-best results are underlined.

Dataset	Component	Prediction					Backtesting			
		ACC	PRE	REC	F1	AUC	AR	SR	CR	MDD
NASDAQ 100	w/o MAGE	<u>52.97</u>	<u>52.90</u>	88.98	66.36	50.29	9.68	0.56	1.09	<u>8.85</u>
	w/o F. 2D Attn	52.17	52.26	95.11	67.46	48.46	<u>13.41</u>	<u>0.80</u>	<u>1.35</u>	9.96
	w/o TCH	52.73	52.56	<u>95.36</u>	<u>67.77</u>	50.97	9.49	0.50	0.92	10.31
	w/o GPH	52.65	52.61	92.16	66.99	47.98	8.44	0.48	0.78	10.76
	MaGNet	53.72	53.09	96.14	68.41	52.24	17.09	1.05	2.09	8.18

indicate that each component contributes critically to MaGNet’s performance. Removing the MAGE block causes the most substantial degradation in trading performance—Annual Return drops from 17.09% to 8.44%, confirming MAGE’s essential role in capturing complex temporal dynamics for both prediction and profitability. The Feature-wise 2D Spatiotemporal Attention and the dual hypergraph components also play important roles, demonstrating that MaGNet’s superior performance arises from the integration of advanced temporal modeling, spatiotemporal feature fusion, and multi-scale relational learning.

5 Conclusion

In this work, we introduced MaGNet, a novel **Mamba dual-hyperGraph Network** designed for stock prediction by synergistically combining advanced temporal modeling with dynamic relational learning. MaGNet integrates the innovative MAGE block for comprehensive temporal-contextual learning, 2D spatiotemporal attention that enable direct feature-to-feature and stock-to-stock modeling while preserving the intrinsic multivariate structure of financial data,

and dual hypergraph framework (TCH and GPH) to dynamically capture temporal-causal relationships and global market structures. Extensive experiments across six major indices validate that MaGNet significantly outperforms state-of-the-art methods in predictive accuracy, investment returns and risk management. Future work could explore incorporating alternative data sources and investigating the interpretability of learned hypergraph structures for market insight generation.

References

- [1] Ayodele Ariyo Adebiyi, Aderemi Oluyinka Adewumi, and Charles Korede Ayo. 2014. Comparison of ARIMA and artificial neural networks models for stock price prediction. *Journal of Applied Mathematics* 2014, 1 (2014), 614342.
- [2] Md Atik Ahamed and Qiang Cheng. 2024. Timemachine: A time series is worth 4 mambas for long-term forecasting. In *ECAI 2024*. IOS Press, 1688–1695.
- [3] A Behrouz, M Santacatterina, and R Zabih. 2024. MambaMixer: Efficient Selective State Space Models with Dual Token and Channel Selection. arXiv:2403.19888 [cs.LG] <https://arxiv.org/abs/2403.19888>
- [4] Kai Chen, Yi Zhou, and Fangyan Dai. 2015. A LSTM-based method for stock returns prediction: A case study of China stock market. In *2015 IEEE international conference on big data (big data)*. IEEE, 2823–2824.

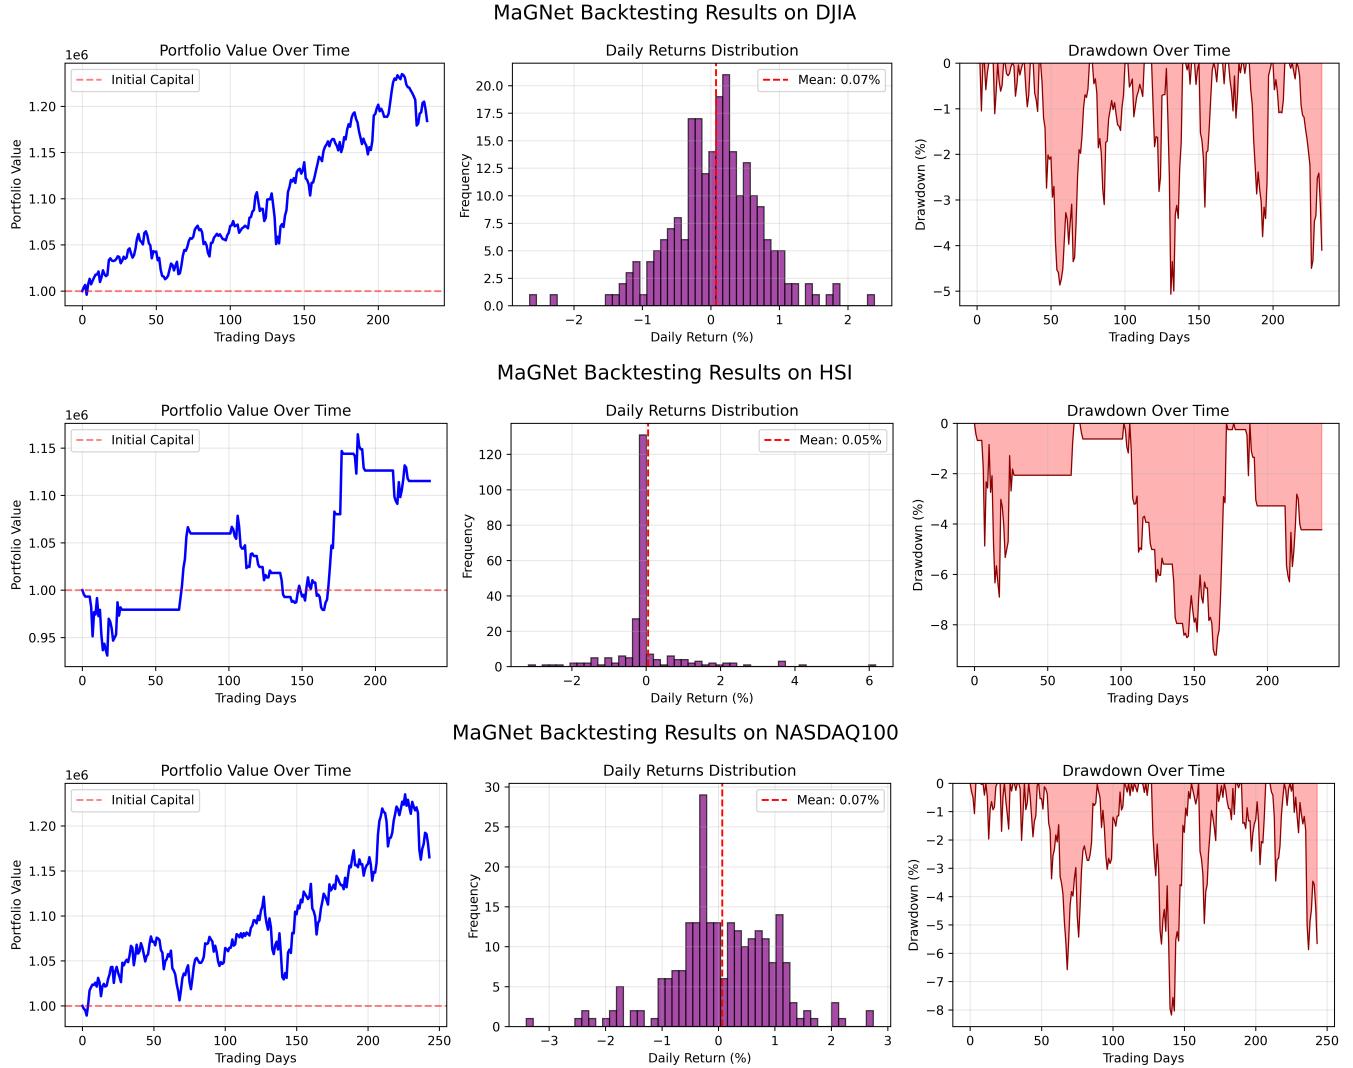


Figure 2: Backtesting performance of the MaGNet on DJIA, HSI and NASDAQ 100 indices

- [5] Weijun Chen, Shun Li, Xipu Yu, Heyuan Wang, Wei Chen, and Tengjiao Wang. 2024. Automatic de-biased temporal-relational modeling for stock investment recommendation. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*. 1999–2008.
- [6] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. 2014. Learning phrase representations using RNN encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078* (2014).
- [7] Chaoran Cui, Xiaojie Li, Chunyun Zhang, Weili Guan, and Meng Wang. 2023. Temporal-relational hypergraph tri-attention networks for stock trend prediction. *Pattern recognition* 143 (2023), 109759.
- [8] Qianggang Ding, Sifan Wu, Hao Sun, Jiadong Guo, and Jian Guo. 2020. Hierarchical multi-scale Gaussian transformer for stock movement prediction. In *Ijcai*. 4640–4646.
- [9] Yitong Duan, Weiran Wang, and Jian Li. 2025. FactorGCL: A Hypergraph-Based Factor Model with Temporal Residual Contrastive Learning for Stock Returns Prediction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. 173–181.
- [10] Jinyong Fan and Yanyan Shen. 2024. StockMixer: A simple yet strong MLP-based architecture for stock price forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 38. 8389–8397.
- [11] Fuli Feng, Huimin Chen, Xiangnan He, Ji Ding, Maosong Sun, and Tat-Seng Chua. 2018. Enhancing stock movement prediction with adversarial training. *arXiv preprint arXiv:1810.09936* (2018).
- [12] Fuli Feng, Xiangnan He, Xiang Wang, Cheng Luo, Yiqun Liu, and Tat-Seng Chua. 2019. Temporal relational ranking for stock prediction. *ACM Transactions on Information Systems (TOIS)* 37, 2 (2019), 1–30.
- [13] Jianliang Gao, Xiaoting Ying, Cong Xu, Jianxin Wang, Shichao Zhang, and Zhao Li. 2021. Graph-based stock recommendation by time-aware relational attention network. *ACM Transactions on Knowledge Discovery from Data (TKDD)* 16, 1 (2021), 1–21.
- [14] Yue Gao, Yifan Feng, Shuyi Ji, and Rongrong Ji. 2022. Hgnn+: General hypergraph neural networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 3 (2022), 3181–3199.
- [15] Felix A Gers, Jürgen Schmidhuber, and Fred Cummins. 2000. Learning to forget: Continual prediction with LSTM. *Neural computation* 12, 10 (2000), 2451–2471.
- [16] Albert Gu and Tri Dao. 2023. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752* (2023).
- [17] Will Hamilton, Zhitao Ying, and Jure Leskovec. 2017. Inductive representation learning on large graphs. *Advances in neural information processing systems* 30 (2017).
- [18] Haodong Han, Liang Xie, Shengshuang Chen, and Haijiao Xu. 2023. Stock trend prediction based on industry relationships driven hypergraph attention networks.

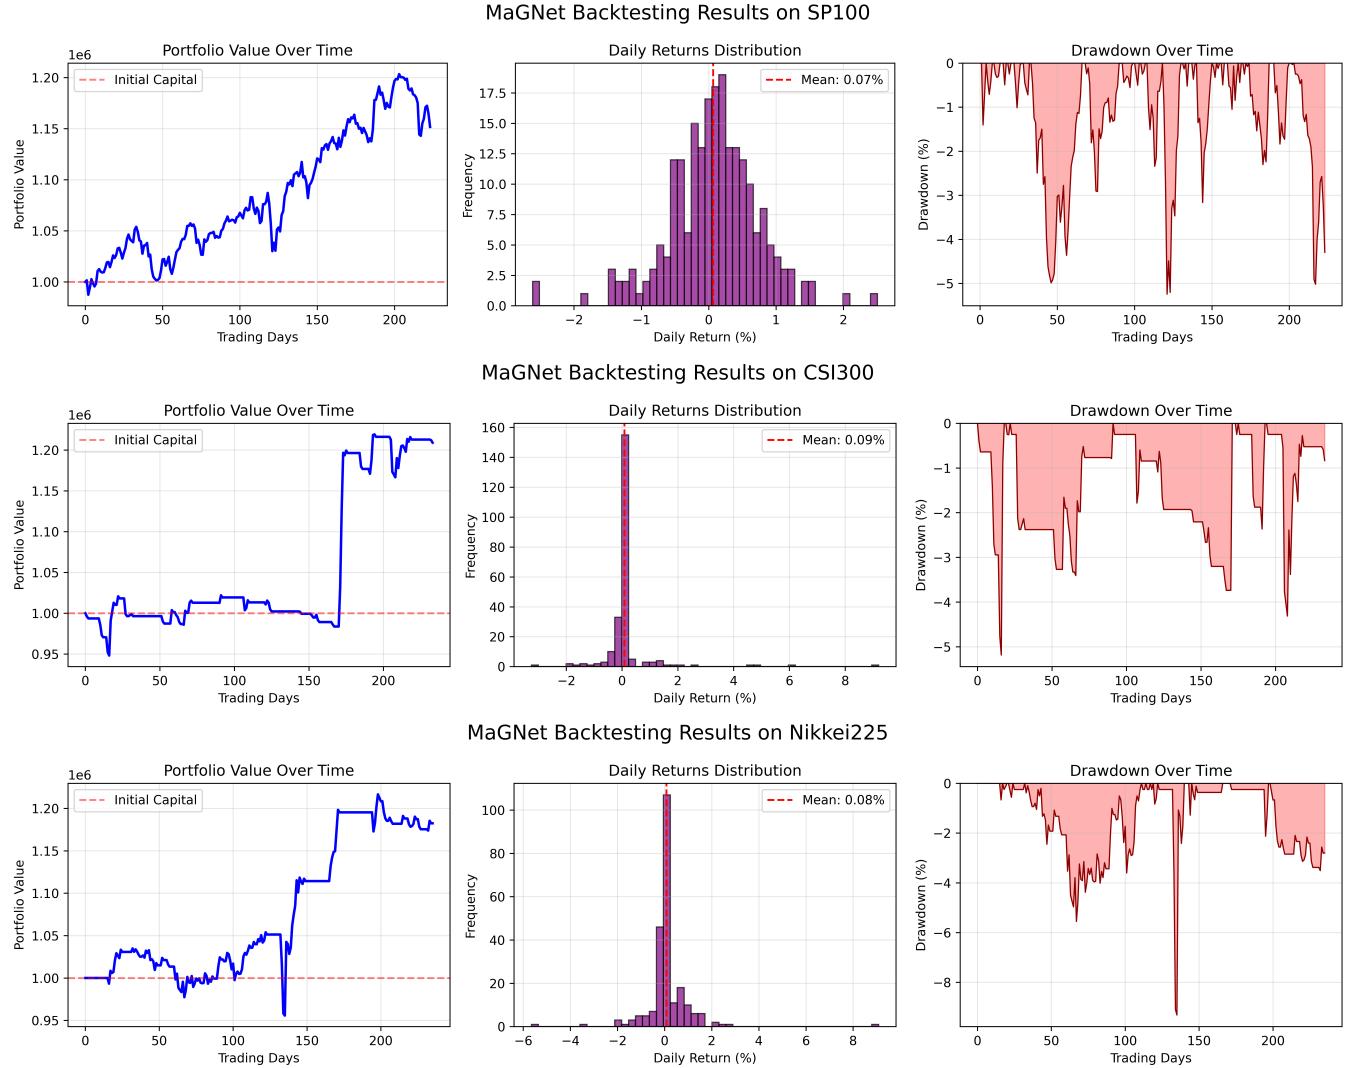


Figure 3: Backtesting performance of the MaGNet on S&P 100, CSI300 and Nikkei225 indices

- Applied Intelligence* 53, 23 (Oct. 2023), 29448–29464. doi:10.1007/s10489-023-05035-z
- [19] Ehsan Hoseinzade and Saman Haratizadeh. 2019. CNNpred: CNN-based stock market prediction using a diverse set of variables. *Expert Systems with Applications* 129 (2019), 273–285.
 - [20] Thanh Trung Huynh, Minh Hieu Nguyen, Thanh Tam Nguyen, Phi Le Nguyen, Matthias Weidlich, Quoc Viet Hung Nguyen, and Karl Aberer. 2023. Efficient integration of multi-order dynamics and internal dynamics in stock movement prediction. In *Proceedings of the Sixteenth ACM International Conference on Web Search and Data Mining*. 850–858.
 - [21] TN Kipf. 2016. Semi-Supervised Classification with Graph Convolutional Networks. *arXiv preprint arXiv:1609.02907* (2016).
 - [22] Hao Li, Yanyan Shen, and Yanmin Zhu. 2018. Stock price prediction using attention-based multi-input LSTM. In *Asian conference on machine learning*. PMLR, 454–469.
 - [23] Tong Li, Zhaoyang Liu, Yanyan Shen, Xue Wang, Haokun Chen, and Sen Huang. 2024. Master: Market-guided stock transformer for stock price forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 162–170.
 - [24] Zhige Li, Derek Yang, Li Zhao, Jiang Bian, Tao Qin, and Tie-Yan Liu. 2019. Individualized indicator for all: Stock-wise technical indicator optimization with stock embedding. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. 894–902.
 - [25] Siyao Liao, Liang Xie, Yuanchuang Du, Shengshuang Chen, Hongyang Wan, and Haijiao Xu. 2024. Stock trend prediction based on dynamic hypergraph spatio-temporal network. *Applied Soft Computing* 154 (2024). 111329.
 - [26] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. 2023. itransformer: Inverted transformers are effective for time series forecasting. *arXiv preprint arXiv:2310.06625* (2023).
 - [27] Ali Mehrabian, Ehsan Hoseinzade, Mahdi Mazloum, and Xiaohong Chen. 2025. Mamba meets financial markets: A graph-mamba approach for stock price prediction. In *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 1–5.
 - [28] Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. 2022. A time series is worth 64 words: Long-term forecasting with transformers. *arXiv preprint arXiv:2211.14730* (2022).
 - [29] Jongho Park, Jaesung Park, Zheyang Xiong, Nayoung Lee, Jaewoong Cho, Samet Oymak, Kangwook Lee, and Dimitris Papailiopoulos. 2024. Can mamba learn how to learn? a comparative study on in-context learning tasks. *arXiv preprint arXiv:2402.04248* (2024).
 - [30] Yao Qin, Dongjin Song, Haifeng Chen, Wei Cheng, Guofei Jiang, and Garrison Cottrell. 2017. A dual-stage attention-based recurrent neural network for time series prediction. *arXiv preprint arXiv:1704.02971* (2017).

- [31] Ramit Sawhney, Shivam Agarwal, Arnav Wadhwa, and Rajiv Ratn Shah. 2020. Spatiotemporal hypergraph convolution network for stock movement forecasting. In *2020 IEEE International Conference on Data Mining (ICDM)*. IEEE, 482–491.
- [32] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. 2017. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. *arXiv preprint arXiv:1701.06538* (2017).
- [33] Shun-Yao Shih, Fan-Keng Sun, and Hung-yi Lee. 2019. Temporal pattern attention for multivariate time series forecasting. *Machine Learning* 108, 8 (2019), 1421–1441.
- [34] Yu-Che Tsai, Chih-Yao Chen, Shao-Lun Ma, Pei-Chi Wang, You-Jia Chen, Yu-Chieh Chang, and Cheng-Te Li. 2019. FineNet: a joint convolutional and recurrent neural network model to forecast and recommend anomalous financial items. In *Proceedings of the 13th ACM conference on recommender systems*. 536–537.
- [35] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).
- [36] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. 2017. Graph attention networks. *arXiv preprint arXiv:1710.10903* (2017).
- [37] Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y Zhang, and Jun Zhou. 2024. Timemixer: Decomposable multiscale mixing for time series forecasting. *arXiv preprint arXiv:2405.14616* (2024).
- [38] Yuxuan Wang, Haixu Wu, Jiaxiang Dong, Guo Qin, Haoran Zhang, Yong Liu, Yunzhong Qiu, Jianmin Wang, and Mingsheng Long. 2024. Timexer: Empowering transformers for time series forecasting with exogenous variables. *Advances in Neural Information Processing Systems* 37 (2024), 469–498.
- [39] Zihan Wang, Fanheng Kong, Shi Feng, Ming Wang, Xiaocui Yang, Han Zhao, Dalong Wang, and Yifei Zhang. 2025. Is mamba effective for time series forecasting? *Neurocomputing* 619 (2025), 129178.
- [40] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. 2022. Timesnet: Temporal 2d-variation modeling for general time series analysis. *arXiv preprint arXiv:2210.02186* (2022).
- [41] Hongjie Xia, Huijie Ao, Long Li, Yu Liu, Sen Liu, Guannan Ye, and Hongfeng Chai. 2024. Ci-sthpam: Pre-trained attention network for stock selection with channel-independent spatio-temporal hypergraph. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 9187–9195.
- [42] Naganand Yadati, Madhav Nimishakavi, Prateek Yadav, Vikram Nitin, Anand Louis, and Partha Talukdar. 2019. Hypergen: A new method for training graph convolutional networks on hypergraphs. *Advances in neural information processing systems* 32 (2019).
- [43] Xiao Yang, Weiqing Liu, Dong Zhou, Jiang Bian, and Tie-Yan Liu. 2020. Qlib: An ai-oriented quantitative investment platform. *arXiv preprint arXiv:2009.11189* (2020).
- [44] Jaemin Yoo, Yejun Soun, Yong-chan Park, and U Kang. 2021. Accurate multivariate stock movement prediction via data-axis transformer with multi-level contexts. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*. 2037–2045.
- [45] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. 2023. Are transformers effective for time series forecasting?. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 37. 11121–11128.
- [46] Liheng Zhang, Charu Aggarwal, and Guo-Jun Qi. 2017. Stock price prediction via discovering multi-frequency trading patterns. In *Proceedings of the 23rd ACM SIGKDD international conference on knowledge discovery and data mining*. 2141–2149.

A Experiment Setting Supplement

A.1 Baseline Descriptions

To evaluate the effectiveness of MaGNet, we compare it against 17 baselines with several state-of-the-art baselines from 3 different categories. These models provide a diverse set of benchmarks to evaluate our method's performance.

1. Stock Prediction Models (6):

- SFM [46]: State Frequency Memory networks that model price fluctuations across multiple frequencies using frequency-based decomposition.
- Adv-ALSTM [11]: Attentive LSTM with adversarial training for improved robustness against stochastic price movements.
- DTML [44]: Transformer architecture capturing dynamic inter-stock correlations through multi-level contexts.

- ESTIMATE [20]: Combines wavelet-based hypergraph convolution with memory-enhanced LSTM for non-pairwise stock correlations.
- StockMixer [10]: MLP-based model that sequentially mixes indicators, temporal patterns, and market correlations.
- MASTER [23]: Integrates intra/inter-stock attention with market-guided gating for dynamic correlation capture.

2. Time Series Models (8):

- GRU [6]: RNN encoder-decoder with gated recurrent units for sequence-to-vector encoding.
- LSTM [15]: RNN architecture with gating mechanisms for long-term dependency modeling.
- DLinear [45]: One-layer linear model that directly models temporal relations for long-term forecasting.
- TimesNet [40]: Transforms time series to 2D tensors to model intra/inter-period variations.
- PatchTST [28]: Channel-independent Transformer using patching for improved long-term forecasting.
- iTransformer [26]: Inverted Transformer applying attention across variates rather than time steps.
- TimeMixer [37]: MLP-based model using multiscale mixing to disentangle temporal variations.
- TimeXer [38]: Transformer designed for forecasting with exogenous variables using patch-wise and variate-wise attention.

3. Graph Models (3):

- GCN (Graph Convolutional Network) [21]: Uses first-order spectral graph convolutions for efficient node embedding learning.
- GraphSAGE [17]: Inductive framework generating embeddings via neighborhood sampling and aggregation.
- GAT (Graph Attention Network) [36]: Employs masked self-attention to assign weights to neighbors for flexible node embedding.

A.2 Metric Definitions

A.2.1 Prediction Metrics.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}, \quad (28)$$

$$\text{Precision} = \frac{TP}{TP + FP}, \quad (29)$$

$$\text{Recall} = \frac{TP}{TP + FN}, \quad (30)$$

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}, \quad (31)$$

$$\text{AUC} = \int_0^1 \text{TPR}(\text{FPR}) d(\text{FPR}), \quad (32)$$

where:

- TP (True Positives): Correctly predicted positive cases;
- TN (True Negatives): Correctly predicted negative cases;
- FP (False Positives): Incorrectly predicted as positive;
- FN (False Negatives): Incorrectly predicted as negative;
- TPR (True Positive Rate) = Recall = $\frac{TP}{TP+FN}$;
- FPR (False Positive Rate) = $\frac{FP}{FP+TN}$.

Table 5: MaGNet’s Hyperparameter Configurations

Dataset	T	# MAGE	# F. 2D Attn	# TCH	# Hyperedges M_1	Top-K	# S. 2D Attn	# GPH	# Hyperedges M_2
DJIA	20	2	1	1	32	32	1	1	16
HSI	10	1	1	2	64	64	2	2	32
NASDAQ 100	10	1	1	2	64	64	1	2	32
S&P100	30	2	1	2	64	64	1	1	32
CSI300	10	1	1	2	64	64	1	1	32
Nikkei 225	10	1	1	1	128	128	2	2	64

A.2.2 Backtesting Metrics.

$$\text{Annual Return} = \left[\prod_{t=1}^T (1 + r_t) \right]^{\frac{252}{T}} - 1, \quad (33)$$

where r_t = return for day t, T = number of trading days, 252 = typical number of trading days per year.

$$\text{Sharpe Ratio} = \frac{R_p - R_f}{\sigma_p}, \quad (34)$$

where R_p = annualized portfolio return, $R_f = 0.02$ (2% risk-free rate), σ_p = annualized standard deviation = $\sigma_{\text{daily}} \times \sqrt{252}$.

$$\text{Calmar Ratio} = \frac{\text{Annual Return}}{|\text{Maximum Drawdown}|}, \quad (35)$$

$$\text{Maximum Drawdown} = \min_{t \in [0, T]} \left(\frac{P_t - \max_{s \in [0, t]} P_s}{\max_{s \in [0, t]} P_s} \right), \quad (36)$$

where P_t = portfolio value at day t.

A.3 Implementation Details

We passed the final representations Z^{GPH} through a FFN followed by softmax to predict the probability distribution of stocks’ close price movement directions (rise/fall) for the next trading day. The model was implemented in PyTorch and trained using cross-entropy loss. Hyperparameters were selected via grid search on each dataset, optimizing for validation accuracy. We tuned the number of layers (1–2) for the MAGE blocks, Feature-wise/Stock-wise Spatiotemporal Attention, TCH and GPH. Key settings included embedding dimension $D = 32$, MoE experts=4, spatiotemporal attention channels=4, attention heads=2 (MAGE and CausalMHA) and dropout=0.1. Lookback windows selected from {10, 20, 30}. Training employed AdamW optimizer with learning rate $1e - 4$. We trained for up to 30 epochs using early stopping. Each layer employed residual connections and layer normalization to support stable deep architecture. All baselines used official implementations with default parameters, adapted for stock prediction using cross-entropy loss. Dataset-specific hyperparameter configurations are shown in Table 5, including Top-K sparsification in TCH; number of hyperedges M_1 for TCH and M_2 for GPH, respectively.

A.4 Backtesing Strategy and Configurations

A.4.1 Dynamic Daily Trading Strategy. In this section, we provide a detailed description of the dynamic daily stock trading strategy and hyperparameter configurations used for backtesting.

Algorithm 1: Dynamic Daily Trading Strategy

Data: N stocks, portfolio proportion p , stop-loss threshold q , conservative ratio r , initial capital 1,000,000, transaction cost rate $\tau = 0.25\%$

```

1 for each trading day  $t$  do
2   // Prediction Generation
3    $P_t \leftarrow \text{Model.predict\_probabilities}(\text{all } N \text{ stocks});$ 
4    $M \leftarrow |\{s : P_{s,t} > 0.5\}|;$ 
   // Portfolio Construction
5    $n_t \leftarrow \begin{cases} \lfloor p \times N \rfloor & \text{if } M \geq p \times N \\ \lfloor r \times M \rfloor & \text{if } p \times N \times q \leq M < p \times N \\ 0 & \text{if } M < p \times N \times q \end{cases}$ 
   // Portfolio Reconstitution & Rebalancing
7   if  $n_t = 0$  then
8     | Liquidate all holdings (apply  $\tau$ );
9   else
10    | Targets $_t \leftarrow \text{Top-}n_t \text{ stocks by } P_t;$ 
11    | Liquidate positions  $\notin$  Targets $_t$  (apply  $\tau$ );
12    // Equal capital allocation
13    | TargetValue  $\leftarrow \frac{\text{TotalPortfolioValue}}{n_t};$ 
14    | for each stock  $s \in \text{Targets}_t$  do
15      | | Adjust position of  $s$  to TargetValue (apply  $\tau$ );
16    | end
17  end
18 end

```

We initialize each backtest with a capital of 1,000,000 and apply a transaction cost rate of 0.25% per trade to emulate realistic market frictions. The trading universe consists of all N stocks in each index. The core design is a dynamic daily trading cycle with adaptive portfolio construction and a stop-loss mechanism. Each trading day proceeds through the following steps. The pseudo code of Dynamic Daily Trading Strategy is shown in Algorithm 1. All transaction adjustment incorporate transaction costs.

- **Prediction Generation:** The model outputs the probability of next-day price rise for all stocks, which we use to rank them.
- **Portfolio Construction with Stop-Loss Mechanism:** We define a portfolio selection proportion p (where $0 < p \leq 1$). On each day:

- If the number of stocks predicted to rise (with probability > 0.5) is at least $p \times N$, we purchase the top $p \times N$ stocks;
- If the number of rising predictions falls into an intermediate zone, specifically $p \times N \times q \leq M < p \times N$ (where q is a stop-loss threshold hyperparameter with $0 < q < 1$), then we adopt a conservative approach: only buy the top $r \times M$ predicted rising stocks (with $0 \leq r \leq 1$);
- If the number of rising predictions is below $p \times N \times q$, we do not buy new positions that day and liquidate all current holdings to avoid downside exposure.
- Portfolio Reconstitution: Positions excluded from the new targets are liquidated with proceeds credited to cash. New target stocks are then purchased with equal capital allocation, subject to current available cash.
- Portfolio Rebalancing: To maintain equal-capital allocations, we adjust positions daily—selling excess holdings that exceed target allocation and purchasing additional shares for under-allocated positions.

This backtesting framework enables direct comparison of model predictions in a realistic trading environment, providing a robust evaluation of each model’s profitability performance under real-world conditions.

Table 6: MaGNet’S Backtesting Hyperparameters

Dataset	P	q	r
DJIA	1	0.05	0
HSI	1	0.05	0
NASDAQ 100	1	0.4	1
S&P 100	1	0.05	0
CSI 300	1	0.05	0
Nikkei 225	1	0.7	1

A.4.2 Backtesting Hyperparameter Configurations. For each (model, dataset) pair, we perform a grid search on the validation set over three parameters in dynamic trading strategy:

- Portfolio Selection Ratio $p \in \{0.05, 0.10, \dots, 1.0\}$;
- Stop-loss Threshold $q \in \{0.05, 0.10, \dots, 0.95\}$;
- Rising Ratio for Partial Entry $r \in \{0.0, 0.05, \dots, 1.0\}$.

The combination yielding the highest Sharpe ratio on the validation set is selected and applied to the test set for final evaluation. The backtesting hyperparameters for MaGNet on each dataset are shown in Table 6.

B Additional Experimental Results

This section presents the additional prediction and backtesting results on the S&P 100, CSI 300, and Nikkei 225 indices, and results of ablation studies on all six indices.

B.1 Additional Prediction Results

Experiment results on S&P 100, CSI 300, and Nikkei 225 indices are shown in Table 7, which further validate the generalizability

and profitability of MaGNet. It achieves the highest accuracy on CSI 300 (54.90%), HSI (54.19%), and Nikkei 225 (54.02%). MaGNet also shows strong recall on S&P 100 (97.00%) and balanced precision-recall trade-offs across datasets. These results confirm that the integration of MAGE’s comprehensive temporal modeling with the dual-hypergraph framework effectively captures both localized temporal-causal dependencies and global market structures, yielding robust and generalizable stock movement predictions across diverse markets, with especially strong improvements on larger universes (CSI 300, Nikkei 225).

B.2 Additional Backtesting Results

Backtesting results om S&P 100, CSI300, and Nikkei 225 are reported in Table 8. On these indices, MaGNet attains the highest Sharpe Ratios of 1.40 (S&P 100), 1.32 (CSI300), and 1.14 (Nikkei 225). It also records the largest Annual Returns among the compared methods on S&P 100 (17.14%), CSI300 (22.60%), and Nikkei 225 (19.58%). Furthermore, MaGNet achieves the highest Calmar ratio across these indices, indicating superior returns with moderate drawdowns (5.19% on CSI 300). These results reinforce the main-text findings that the MAGE block’s temporal modeling combined with the dual-hypergraph architecture yields consistently higher risk-adjusted returns and competitive downside control, particularly on larger market universes.

B.3 Complete Results of Ablation Studies

Comprehensive ablation results on all six indices are reported in Table 9. The results reveal that each component contributes critically to MaGNet’s performance, with varying impacts across different markets.

Removing the MAGE block causes the most substantial degradation, particularly in trading performance—Annual Return on DJIA drops from 19.92% to 9.38% and Sharpe Ratio from 1.70 to 0.70, while recall on CSI 300 plummets from 34.58% to 11.90%, confirming MAGE block’s essential role in capturing complex temporal dynamics for both prediction and profitability.

The Feature-wise 2D Spatiotemporal Attention proves crucial for maintaining stable trading performance, as Annual Return on S&P 100 drops from 17.14% to 13.61% without it, validating its effectiveness in preserving cross-feature spatiotemporal structures.

The dual-hypergraph components exhibit complementary and market-specific contributions: TCH removal significantly impacts markets with strong temporal dependencies (AR on HSI drops from 12.25% to 8.59%), while GPH removal substantially degrades performance in markets with prominent global patterns (AR on NASDAQ 100 falls from 17.09% to 8.44%).

Overall, the ablation studies demonstrate that MaGNet’s superior performance stems from the synergistic integration of advanced temporal modeling through MAGE block, 2D spatiotemporal feature fusion, and multi-scale relational learning via the dual-hypergraph framework, with each module addressing distinct aspects of the complex stock prediction challenge.

Table 7: Prediction performance comparison on S&P 100, CSI 300, and Nikkei 225. The best results are in bold and the second-best results are underlined.

Model	S&P 100					CSI 300					Nikkei 225				
	ACC	PRE	REC	F1	AUC	ACC	PRE	REC	F1	AUC	ACC	PRE	REC	F1	AUC
GRU	51.57	52.75	76.71	62.52	50.46	53.00	52.85	10.86	18.02	52.26	51.28	52.85	48.61	50.64	52.11
LSTM	52.37	52.76	91.07	66.81	50.76	53.06	53.17	10.79	17.94	52.88	50.12	52.04	38.14	44.02	50.92
DLinear	52.29	52.61	94.12	67.50	<u>51.21</u>	52.97	51.28	22.25	31.04	52.43	51.18	51.36	95.00	66.68	49.94
TimesNet	52.26	52.87	85.72	65.40	50.42	53.02	50.91	33.94	40.73	53.22	<u>51.66</u>	52.82	56.13	54.43	52.46
PatchTST	50.78	52.64	64.71	58.05	49.93	51.86	49.18	37.29	42.42	51.77	49.31	50.66	54.29	52.41	48.94
iTransformer	52.32	52.65	93.56	67.38	50.04	52.27	49.65	26.46	34.52	51.95	49.09	50.55	45.69	48.00	48.92
TimeMixer	52.54	52.69	96.42	68.14	50.28	52.78	52.19	8.42	14.50	52.99	51.42	51.42	100.00	67.92	48.38
TimeXer	52.48	52.66	96.24	68.07	50.27	52.98	52.52	11.58	18.97	52.77	51.42	51.42	100.00	67.92	51.15
GCN	50.31	52.03	71.81	60.34	50.80	<u>54.13</u>	<u>53.50</u>	27.04	35.92	52.33	50.97	52.64	46.18	49.20	51.26
GraphSAGE	52.16	52.66	90.29	66.52	49.78	53.06	52.24	14.96	23.26	52.05	51.33	<u>52.88</u>	49.05	50.89	52.91
GAT	49.77	52.41	49.71	51.02	50.09	50.64	47.90	43.46	45.57	51.51	51.03	52.12	58.48	55.12	<u>53.03</u>
SFM	50.76	53.84	52.50	53.16	50.95	53.12	50.78	<u>39.74</u>	<u>44.59</u>	<u>53.95</u>	50.72	52.27	45.23	48.50	51.05
Adv-ALSTM	52.45	52.65	94.91	67.71	50.01	53.00	51.82	17.49	26.05	52.89	51.41	52.59	56.55	54.44	51.81
DTML	52.44	52.88	88.68	66.25	51.03	53.38	52.74	18.83	27.75	53.59	51.05	52.19	57.10	54.54	51.51
ESTIMATE	<u>52.59</u>	52.71	<u>96.70</u>	<u>68.23</u>	50.25	52.37	46.64	1.11	2.16	49.40	50.22	51.83	45.23	48.31	50.43
StockMixer	52.53	52.69	95.99	68.04	50.98	52.77	50.52	32.93	39.87	53.24	51.40	51.47	96.11	67.04	50.47
MASTER	51.66	<u>53.72</u>	58.97	56.22	51.68	52.94	51.05	25.07	33.63	53.13	51.60	51.56	97.20	67.38	52.47
MaGNet	53.14	53.00	97.00	68.55	49.36	54.90	54.03	34.58	42.17	54.59	54.02	54.96	58.59	56.72	53.55

Table 8: Backtesting performance comparison on S&P 100, CSI300, and Nikkei 225. The best results are in bold and the second-best results are underlined.

Model	S&P 100				CSI 300				Nikkei 225			
	AR	SR	CR	MDD	AR	SR	CR	MDD	AR	SR	CR	MDD
GRU	13.59	1.12	1.98	6.85	16.35	<u>1.27</u>	1.83	8.93	9.97	0.52	0.92	10.80
LSTM	13.70	1.06	2.03	6.74	20.22	1.26	2.79	7.25	14.23	0.47	0.66	21.61
DLinear	10.74	0.80	1.84	5.84	16.61	1.15	1.97	8.43	7.13	0.22	0.31	22.91
TimesNet	7.00	0.45	1.06	6.58	7.50	0.55	0.77	9.74	3.40	0.07	0.21	15.94
PatchTST	1.49	-0.15	0.83	<u>1.80</u>	0.81	-0.87	1.04	0.78	3.18	0.10	0.49	6.47
iTransformer	11.81	0.90	2.15	5.50	12.74	0.74	2.21	5.77	5.17	0.32	1.28	4.04
TimeMixer	14.87	1.19	2.75	5.40	<u>21.38</u>	0.90	2.24	9.55	15.18	0.57	0.67	22.57
TimeXer	14.27	1.13	2.55	5.60	9.03	0.61	1.37	6.60	15.18	0.57	0.67	22.57
GCN	<u>16.99</u>	0.91	1.58	10.79	14.55	1.18	2.27	6.41	15.33	<u>0.81</u>	1.04	14.78
GraphSAGE	<u>10.24</u>	0.79	1.84	5.56	12.47	0.77	1.33	9.35	13.96	<u>0.56</u>	0.66	21.05
GAT	9.87	0.69	1.44	6.84	9.89	0.44	0.60	16.40	16.39	0.57	0.67	24.42
SFM	1.98	-0.01	1.45	1.36	14.61	0.98	2.09	7.00	6.43	0.76	<u>1.40</u>	<u>4.59</u>
Adv-ALSTM	15.65	<u>1.27</u>	<u>2.97</u>	5.27	16.28	0.84	1.13	14.35	5.69	0.20	0.32	17.57
DTML	12.69	0.97	2.12	5.98	14.00	1.04	<u>2.47</u>	5.67	6.21	0.33	0.60	10.37
ESTIMATE	15.58	1.16	2.93	5.33	15.37	0.59	1.58	9.76	<u>19.57</u>	0.66	0.75	26.17
StockMixer	11.83	0.92	2.09	5.66	12.59	1.11	2.11	5.97	9.32	0.32	0.41	22.95
MASTER	4.10	0.35	1.08	3.78	15.65	1.15	2.39	6.54	17.58	0.69	0.78	22.58
MaGNet	17.14	1.40	3.27	5.24	22.60	1.32	4.36	<u>5.19</u>	19.58	1.14	2.10	9.31

Table 9: Complete ablation results. The best results are in bold and the second-best results are underlined.

Dataset	Component	Prediction					Backtesting			
		ACC	PRE	REC	F1	AUC	AR	SR	CR	MDD
DJIA	w/o MAGE	52.73	53.13	94.13	67.92	49.61	9.38	0.70	1.34	7.02
	w/o F. 2D Attn	<u>53.10</u>	<u>53.15</u>	<u>99.30</u>	<u>69.24</u>	<u>50.84</u>	15.61	1.30	2.83	5.51
	w/o TCH	53.03	53.12	99.19	69.19	49.68	<u>18.32</u>	<u>1.54</u>	<u>3.62</u>	5.07
	w/o GPH	52.96	53.08	99.19	69.16	48.73	<u>18.16</u>	1.53	3.38	5.37
	MaGNet	53.16	53.16	100.00	69.42	51.10	19.92	1.70	3.93	5.07
HSI	w/o MAGE	52.43	51.15	24.01	32.68	50.00	6.98	0.35	0.49	14.23
	w/o F. 2D Attn	52.34	50.80	<u>28.09</u>	<u>36.18</u>	50.95	6.44	0.31	<u>0.67</u>	<u>9.61</u>
	w/o TCH	52.53	52.30	14.49	22.69	50.44	<u>8.59</u>	0.29	0.39	21.80
	w/o GPH	<u>53.76</u>	54.07	25.42	34.59	<u>53.60</u>	8.26	<u>0.36</u>	0.59	13.89
	MaGNet	54.19	52.85	43.90	47.96	54.12	12.25	0.66	1.33	9.20
NASDAQ 100	w/o MAGE	<u>52.97</u>	<u>52.90</u>	88.98	66.36	50.29	9.68	0.56	1.09	<u>8.85</u>
	w/o F. 2D Attn	52.17	52.26	95.11	67.46	48.46	<u>13.41</u>	<u>0.80</u>	<u>1.35</u>	9.96
	w/o TCH	52.73	52.56	<u>95.36</u>	<u>67.77</u>	<u>50.97</u>	9.49	0.50	0.92	10.31
	w/o GPH	52.65	52.61	92.16	66.99	47.98	8.44	0.48	0.78	10.76
	MaGNet	53.72	53.09	96.14	68.41	52.24	17.09	1.05	2.09	8.18
S &P 100	w/o MAGE	52.74	52.89	93.63	67.59	50.14	10.63	0.81	1.93	5.52
	w/o F. 2D Attn	<u>52.97</u>	<u>52.95</u>	95.84	<u>68.21</u>	49.23	<u>13.61</u>	<u>1.13</u>	<u>2.59</u>	5.24
	w/o TCH	52.10	52.46	95.88	67.82	48.81	<u>12.97</u>	1.01	2.08	6.23
	w/o GPH	52.42	52.61	<u>96.84</u>	68.18	48.98	12.49	0.97	2.38	5.24
	MaGNet	53.14	53.00	97.00	68.55	<u>49.36</u>	17.14	1.40	3.27	5.24
CSI 300	w/o MAGE	<u>54.40</u>	60.46	11.90	19.89	51.40	8.10	0.57	1.48	<u>5.45</u>
	w/o F. 2D Attn	53.07	52.29	14.89	23.17	50.93	<u>18.36</u>	<u>0.94</u>	1.48	12.42
	w/o TCH	52.91	51.04	23.83	32.49	<u>52.68</u>	12.12	0.66	<u>2.06</u>	5.88
	w/o GPH	53.25	51.43	<u>30.26</u>	<u>38.10</u>	51.67	11.84	0.62	1.13	10.47
	MaGNet	54.90	54.03	34.58	42.17	54.59	22.60	1.32	4.36	5.19
Nikkei 225	w/o MAGE	52.74	<u>54.60</u>	48.10	51.14	<u>53.38</u>	3.26	0.07	0.18	18.56
	w/o F. 2D Attn	53.05	54.47	53.00	53.73	<u>52.69</u>	<u>14.25</u>	<u>0.59</u>	0.63	22.57
	w/o TCH	51.01	52.34	52.77	52.56	50.98	9.24	0.55	<u>1.05</u>	8.84
	w/o GPH	<u>53.13</u>	54.15	<u>57.67</u>	<u>55.86</u>	51.97	2.14	0.01	0.09	22.94
	MaGNet	54.02	54.96	58.59	56.72	53.55	19.58	1.14	2.10	<u>9.31</u>